

World Model

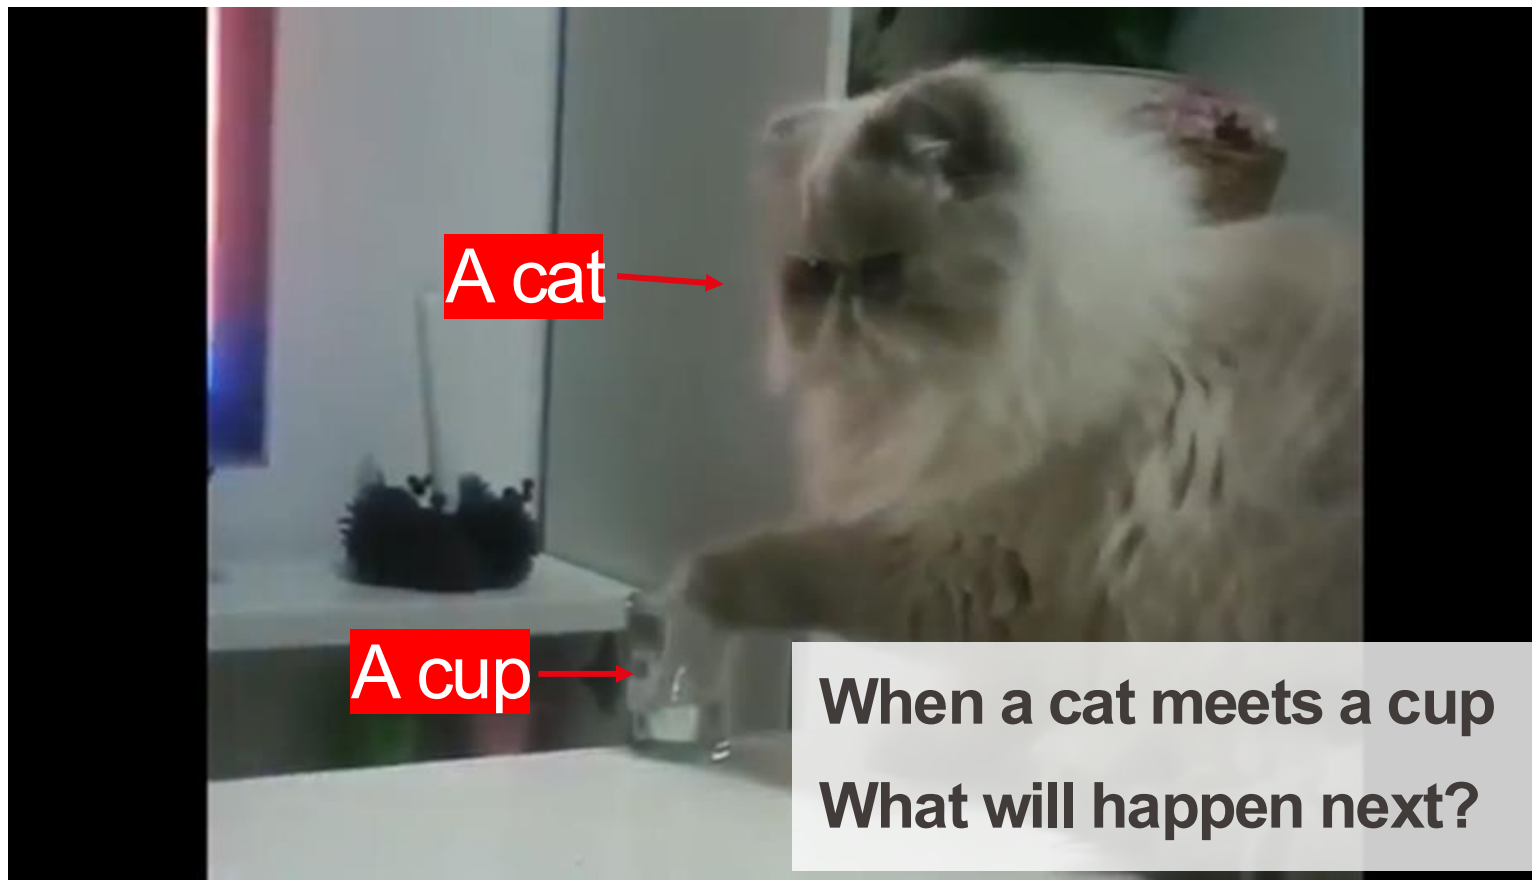
The Intelligent Brain of the Next-Generation Autonomous Systems

Presented By: **Mohamed** Abdelfattah (Os)


Slides by **Wuyang** Li

April 23, 2026

What is World Model?



What is World Model?




We need a new one!



EPFL What is World Model?



When you see the **current scene** and potential **action**



You can **imagine** the future

What is World Model?

Why we have such ability ?

EPFL What is World Model?

Why we have such ability ?


Because our brain secretly has a “mental world”



EPFL What is World Model?

Why we have such ability ?

Because our brain secretly has a “mental model”



EPFL What is World Model?

Why we have such ability ?


Predict the **Future** given the
Current **Observation** and **Action**




1 Delving into World Model




- Humans learn best by **actively engaging** with the world.




- Humans learn best by **actively engaging** with the world.
 - It is challenging for robots




- Humans learn best by actively engaging with the world.
 - It is challenging for robots
- How about **simulating a world** for robots?





With World Model, We CAN...



With World Model, We CAN...



1 Delving into World Model

2 Classic World Model




- Developing an **autonomous player** to achieve the goal in game



- Dreamer [1]

- Explicit model **predicting the future image** and reward



- TD-MPC [2] (Temporal Difference Learning for Model Predictive Control)

- Explicit model **predicting the future state** and reward



- MuZero [3]

- Implicit model only making decisions


[1] Hafner et al. Dream to control: Learning behaviors by latent imagination. ICLR 2020

[2] Hansen et al. Temporal difference learning for model predictive control. ICML 2022

[3] Schrittwieser et al. Mastering atari, go, chess and shogi by planning with a learned model. Nature 2020.

Dreamer: Autonomous Player

- DreamerV3 is the first algorithm that can collect diamonds in Minecraft without human data (Nature 2025)



Dreamer: Autonomous Player

Step1: World Learning

Train a **world model** as brain to predict the future



Step2: Policy Learning

Based on the world model, learn policy with **reinforcement learning**



Dreamer: World Learning

Step1: World Learning


Train a **world model** as brain to predict the future



Step2: Policy Learning

Based on the world model, learn policy with **reinforcement learning**







- Encoder & Decoder

- Encode images into latent

$$z_t \sim q_\phi(z_t | h_t, x_t)$$

- Decode images from latent

$$\hat{x}_t \sim p_\phi(\hat{x}_t | h_t, z_t)$$



- Encoder & Decoder

- Encode images into latent

$$z_t \sim q_\phi(z_t | h_t, x_t)$$


- Decode images from latent

$$\hat{x}_t \sim p_\phi(\hat{x}_t | h_t, z_t)$$

- Sequential Model

- Predict the future state

$$h_t = f_\phi(h_{t-1}, z_{t-1}, a_{t-1})$$



- Encoder & Decoder

- Encode images into latent

$$z_t \sim q_\phi(z_t | h_t, x_t)$$

- Decode images from latent

$$\hat{x}_t \sim p_\phi(\hat{x}_t | h_t, z_t)$$


- Sequential Model

- Predict the future state

$$h_t = f_\phi(h_{t-1}, z_{t-1}, a_{t-1})$$

- Predictors


- Reward predictor $\hat{z}_t \sim p_\phi(\hat{z}_t | h_t)$
- Continue predictor $\hat{c}_t \sim p_\phi(\hat{c}_t | h_t, z_t)$



Dreamer: Policy Learning

Step1: World Learning

Train a **world model** as brain to predict the future





Step2: Policy Learning

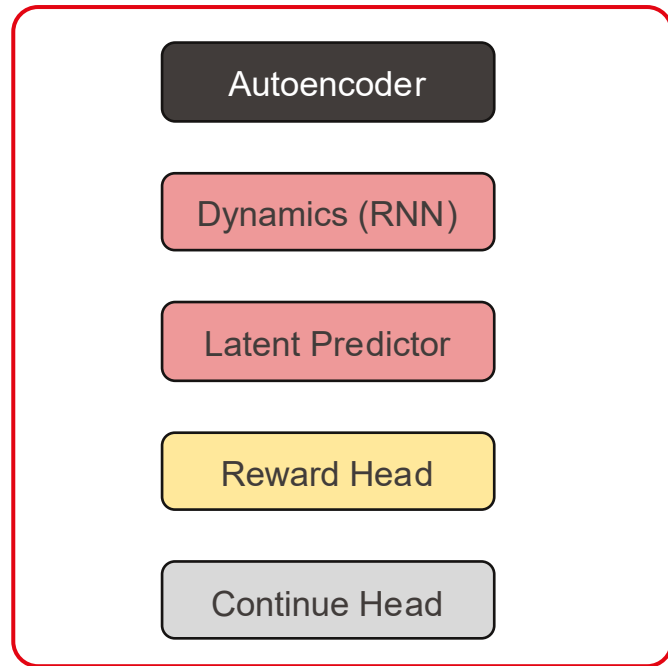
Based on the world model, learn policy with **reinforcement learning**




- Based on the trained world model
 - Predict the future **action**
 - Predict the **value** of future state
- Optimize via reinforcement learning



World Model





Policy



EPFL TD-MPC: Decoder-free World Model

- Finish everything in the latent space
 - **Decoder-free** architecture
 - Directly predict the actions



- Key Components in World Model
 - Autoencoder 
 - Sequential Model 
 - Reward Predictor (optional) 
- Three types of World Model
 - Dreamer
 - TD-MPC
 - MuZero

1 Delving into World Model


2 Classic World Model

3 World Model for Driving



Why We Need World Model

Driving requires decision-making at every moment



Current observations




Actions

Why We Need World Model

We need to know what will happen if we...



Generating Future given Driving Actions





EPFL World Model for Autonomous Driving

35


delfattah

- **What** to generate
 - RGB Videos
 - Depth
 - Lidar
 - Occupancy / BEV
- **How** to generate
 - Diffusion Model
 - Autoregressive Transformer
 - Mixed



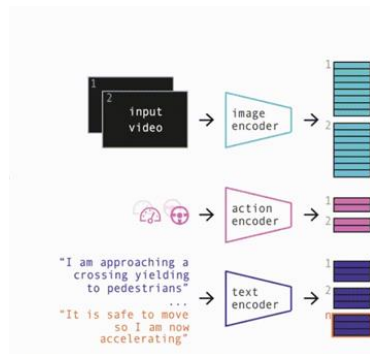
EPFL GAIA: The First Driving World Model based on Video Generation






- Encode into discrete token
- Recurrently generate future token
- Decode token to videos

Hu et al. Gaia-1: A generative world model for autonomous driving. arXiv preprint arXiv:2309.17080 (2023).



Discrete!




- Vector Quantized Variational Autoencoder (VQ-VAE)
 - Encode → Look Up & Replace → Decode

Hu et al. Gaia-1: A generative world model for autonomous driving. arXiv preprint arXiv:2309.17080 (2023).

Van et al. Neural discrete representation learning. NeurIPS 2017

GAIA: Autoregressive Prediction




- GPT-like Autoregressive Prediction
 - Predicting the next token with previous tokens

$$p(\mathbf{x}) = p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

Hu et al. Gaia-1: A generative world model for autonomous driving. arXiv preprint arXiv:2309.17080 (2023).

GAIA: Autoregressive Prediction





- GPT-like Autoregressive Prediction
 - Predicting t-th token with previous tokens

$$p(\mathbf{x}) = p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

Hu et al. Gaia-1: A generative world model for autonomous driving. arXiv preprint arXiv:2309.17080 (2023).

EPFL GAIA: The First Driving World Model



- Video Diffusion Decoder
 - Generate videos with diffusion


Hu et al. "Gaia-1: A generative world model for autonomous driving." arXiv preprint arXiv:2309.17080 (2023).

EPFL Preliminaries: Video Diffusion Model

42

Mohamed Abdelfattah

- Generate videos with gradual denoising
 - Architecture: 3D U-Net



GAIA: The First Driving World Model

- Context & Short Text Control



Text:
"It's raining"



EPFL GAIA: The First Driving World Model

- Limitation in GAIA
 - Low-resolution (288×512)
 - Limited controls (text)



EPFL Vista: Higher Resolution & More Controls

45


Mohamed Abdelfattah



EPFL Vista: Two Technical Improvements

46

Abdelfattah



- Left: **Latent** video diffusion model
- Right: Inject **visual** controls (e.g., trajectory, target point etc.)

Rombach et al. High-resolution image synthesis with latent diffusion models. *CVPR 2022*

- Limitation in Vista
 - Low-dimensional generation lacking **3D understanding**
 - Limited controls of **environment**

2D video



Large Gap
←→

3D scene



EPFL GEM: Multi-modal Future Generation

48


Mohamed Abdelfattah



■ World Model for Driving

Hassan et al. GEM: A Generalizable Ego-Vision Multimodal World Model for Fine-Grained Ego-Motion, Object Dynamics, and Scene Composition Control. CVPR 2025


EPFL GEM: Two Technical Improvements



- Generate videos and corresponding **depth**
- Control the **environment** with diverse visual guidance


Hassan et al. GEM: A Generalizable Ego-Vision Multimodal World Model for Fine-Grained Ego-Motion, Object Dynamics, and Scene Composition Control. CVPR 2025

- Human behaviors
 - E.g., pedestrians cross the road



- Hassan et al. GEM: A Generalizable Ego-Vision Multimodal World Model for Fine-Grained Ego-Motion, Object Dynamics, and Scene Composition Control. CVPR 2025



- Actions of other objects
 - E.g., a car appearing in the blue area



- Hassan et al. GEM: A Generalizable Ego-Vision Multimodal World Model for Fine-Grained Ego-Motion, Object Dynamics, and Scene Composition Control. CVPR 2025


More Interesting World Models

- Generate 3D occupancy (Voxels)



More Interesting World Models

- Generate bird's-eye-view maps



- Control the generation with vectorized maps



EPFL COSMOS: Cutting-Edge World Model

55

Mohamed Abdelfattah







EPFL COSMOS: Cutting-Edge World Model

56

Mohamed Abdelfattah

Input

The video is captured from a camera mounted on a car. The camera is facing forward. The footage shows a quiet American street corner on a sunny day. A small diner with a retro sign stands at one corner, while the other corners feature low-rise office buildings and wide sidewalks. A single pickup truck waits at the intersection, with its brake lights visible with no visible movement. Suddenly, the car collides with the back of the truck. The video is captured from a camera mounted on a car. The camera is facing to the left. The footage shows a wide sidewalk with evenly spaced trees casting shadows on the pavement. A bicycle rack outside a coffee shop holds a single bike, and the shop's large windows glisten in the sunlight. The scene feels open and calm, with no visible movement. The video is captured from a camera mounted on a car. The camera is facing to the right. The footage captures a parking lot with a few scattered vehicles near a convenience store. The bright sunlight casts clear shadows on the asphalt, and the quietness of the street enhances the laid-back atmosphere of the corner. The video is captured from a camera mounted on a car. The camera is facing backwards. The footage shows the road stretching away from the intersection, lined by single-story buildings and an occasional parked car. The sunlight brightens the muted tones of the scene, creating a relaxed and spacious feel. The video is captured from a camera mounted on a car. The camera is facing the rear left side. The video is captured from a camera mounted on a car. The camera is facing the rear right side.

**Output****Input****Segmentation****Depth****Vis****Edge****Output****Input****Output**

The video shows a night-time driving scene where a vehicle is about to make a left turn

- 1 Delving into World Model
- 2 Classic World Model
- 3 World Model for Driving

4 Takeaway Message



- World Model
 - Predict the **future** given the current **observation** and **action**
- Basic Components
 - Autoencoder
 - Sequential Model
 - Predictor
- Basic Generative Pipeline
 - Diffusion Model
 - Autoregressive Model
 - Mixed



<https://www.youtube.com/watch?v=tx3DgeqHO5s>

Thank you for your attention!

VITA Lab

